

k-NN Classification Model

Melissa Paniagua

9/26/2020

The purpose of this assignment is to use k-NN for classification utilizing open data on 5000 customers from a financial institution to predict whether a liability customer would accept a personal loan offer.

Variables used in the Universal Bank dataset are the following:

- ID: Customer's identifier.
- Age: Customer's age.
- Experience: Number of customer's years experience.
- Income: Annual income.
- Zip Code: Customer's location area.
- Family: Number of family members.
- CCAvg: Average spending on credit cards.
- Education: Highest customer's education level.

1 = High School 2 = Undergraduate. 3 = Graduate.

- Mortgage: Value of debt if the customer has a mortgage.
- Personal Loan: Indicates if the customer accepted or rejected the loan offered in last campaign.

1 = Accepted 0 = Rejected

- Securities Account: Indicates if the customer has security account.

1 = Yes. 0 = No.

- CD Account: Indicates if the customer has a Certificate of Deposit.

1 = Yes. 0 = No.

- Online: Indicates if the customer has internet banking facilities.

1 = Yes. 0 = No.

- CreditCard: Indicates if the customer currently has credit cards.

1 = Yes. 0 = No.

```

# Load the libraries needed for the project
library(readr) #read files
library(dplyr) #To select a subset of variables
library(dummies) #To create dummies variables
library(fastDummies) #To create dummies variables
library(caret) #To split the dataset in training, validation, and testing.
library(FNN) #To use k-NN algorithm
library(ggplot2) #To create plots
library("gmodels") # To create the confusion matrix

#Import the dataset
df <- read_csv("UniversalBank.csv")

```

```

#Preview the data
head(df)

```

```

# A tibble: 6 x 14
  ID   Age Experience Income 'ZIP Code' Family CCAvg Education Mortgage
<dbl> <dbl>      <dbl>  <dbl>      <dbl>  <dbl> <dbl>      <dbl>      <dbl>
1     1    25         1     49      91107     4   1.6         1         0
2     2    45        19     34      90089     3   1.5         1         0
3     3    39        15     11      94720     1   1         1         0
4     4    35         9    100      94112     1   2.7         2         0
5     5    35         8     45      91330     4   1         2         0
6     6    37        13     29      92121     4   0.4         2        155
# ... with 5 more variables: 'Personal Loan' <dbl>, 'Securities Account' <dbl>,
#   'CD Account' <dbl>, Online <dbl>, CreditCard <dbl>

```

Exploration Data

```

# See the dataframe's structure
str(df)

```

```

Classes 'spec_tbl_df', 'tbl_df', 'tbl' and 'data.frame':   5000 obs. of  14 variables:
 $ ID          : num  1 2 3 4 5 6 7 8 9 10 ...
 $ Age         : num  25 45 39 35 35 37 53 50 35 34 ...
 $ Experience   : num  1 19 15 9 8 13 27 24 10 9 ...
 $ Income       : num  49 34 11 100 45 29 72 22 81 180 ...
 $ ZIP Code    : num  91107 90089 94720 94112 91330 ...
 $ Family       : num  4 3 1 1 4 4 2 1 3 1 ...
 $ CCAvg        : num  1.6 1.5 1 2.7 1 0.4 1.5 0.3 0.6 8.9 ...
 $ Education    : num  1 1 1 2 2 2 2 3 2 3 ...
 $ Mortgage     : num  0 0 0 0 0 155 0 0 104 0 ...
 $ Personal Loan : num  0 0 0 0 0 0 0 0 0 1 ...
 $ Securities Account: num  1 1 0 0 0 0 0 0 0 0 ...
 $ CD Account   : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Online       : num  0 0 0 0 0 1 1 0 1 0 ...
 $ CreditCard   : num  0 0 0 0 1 0 0 1 0 0 ...
- attr(*, "spec")=
 .. cols(
 ..   ID = col_double(),

```

```

.. Age = col_double(),
.. Experience = col_double(),
.. Income = col_double(),
.. 'ZIP Code' = col_double(),
.. Family = col_double(),
.. CCAvg = col_double(),
.. Education = col_double(),
.. Mortgage = col_double(),
.. 'Personal Loan' = col_double(),
.. 'Securities Account' = col_double(),
.. 'CD Account' = col_double(),
.. Online = col_double(),
.. CreditCard = col_double()
.. )

```

```

#Some descriptive statistics
summary(df)

```

```

      ID      Age      Experience      Income      ZIP Code
Min.   : 1  Min.   :23.00  Min.   : -3.0  Min.   : 8.00  Min.   : 9307
1st Qu.:1251 1st Qu.:35.00 1st Qu.:10.0 1st Qu.: 39.00 1st Qu.:91911
Median :2500 Median :45.00 Median :20.0 Median : 64.00 Median :93437
Mean   :2500 Mean   :45.34 Mean   :20.1 Mean   : 73.77 Mean   :93152
3rd Qu.:3750 3rd Qu.:55.00 3rd Qu.:30.0 3rd Qu.: 98.00 3rd Qu.:94608
Max.   :5000 Max.   :67.00 Max.   :43.0 Max.   :224.00 Max.   :96651

      Family      CCAvg      Education      Mortgage
Min.   :1.000  Min.   : 0.000  Min.   :1.000  Min.   : 0.0
1st Qu.:1.000 1st Qu.: 0.700 1st Qu.:1.000 1st Qu.: 0.0
Median :2.000 Median : 1.500 Median :2.000 Median : 0.0
Mean   :2.396 Mean   : 1.938 Mean   :1.881 Mean   : 56.5
3rd Qu.:3.000 3rd Qu.: 2.500 3rd Qu.:3.000 3rd Qu.:101.0
Max.   :4.000 Max.   :10.000 Max.   :3.000 Max.   :635.0

Personal Loan  Securities Account  CD Account      Online
Min.   :0.000  Min.   :0.0000  Min.   :0.0000  Min.   :0.0000
1st Qu.:0.000 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000
Median :0.000 Median :0.0000 Median :0.0000 Median :1.0000
Mean   :0.096 Mean   :0.1044 Mean   :0.0604 Mean   :0.5968
3rd Qu.:0.000 3rd Qu.:0.0000 3rd Qu.:0.0000 3rd Qu.:1.0000
Max.   :1.000 Max.   :1.0000 Max.   :1.0000 Max.   :1.0000

      CreditCard
Min.   :0.000
1st Qu.:0.000
Median :0.000
Mean   :0.294
3rd Qu.:1.000
Max.   :1.000

```

```

# Identify unique values on specific variables, which could be change to dummies variables.
unique(df$Education)

```

```
[1] 1 2 3
```

Dummies variables

It is important to transform categorical variables and transform them into dummies variables. The column we will transform is education.

```
# Create dummies variables for education.
dumdf <- dummy_cols(df, select_columns = c("Education"))

# See the dataframe with dummies variables
head(dumdf)
```

```
# A tibble: 6 x 17
  ID Age Experience Income 'ZIP Code' Family CCAvg Education Mortgage
  <dbl> <dbl>      <dbl> <dbl>      <dbl> <dbl> <dbl>      <dbl>      <dbl>
1     1    25         1     49      91107     4    1.6         1         0
2     2    45        19     34      90089     3    1.5         1         0
3     3    39        15     11      94720     1     1         1         0
4     4    35         9    100      94112     1    2.7         2         0
5     5    35         8     45      91330     4     1         2         0
6     6    37        13     29      92121     4    0.4         2        155
# ... with 8 more variables: 'Personal Loan' <dbl>, 'Securities Account' <dbl>,
# 'CD Account' <dbl>, Online <dbl>, CreditCard <dbl>, Education_1 <int>,
# Education_2 <int>, Education_3 <int>
```

Select a subset of variables

For our project purposes, we will create a new dataset ignoring variables that are not important for our project such as ID number, zip code information, and also the original education columns.

```
#Select subset.
mydf <- select(dumdf, 2:4, 6:7, 9:17)

# Move the predicted variable (Personal Loan) at the end of the dataset.
mydf <- mydf%>%select(-'Personal Loan', 'Personal Loan')

#See new dataframe
head(mydf)
```

```
# A tibble: 6 x 14
  Age Experience Income Family CCAvg Mortgage 'Securities Acc~ 'CD Account'
  <dbl>      <dbl> <dbl> <dbl> <dbl>      <dbl>      <dbl>      <dbl>
1    25         1     49     4    1.6         0         1         0
2    45        19     34     3    1.5         0         1         0
3    39        15     11     1     1         0         0         0
4    35         9    100     1    2.7         0         0         0
5    35         8     45     4     1         0         0         0
6    37        13     29     4    0.4        155         0         0
# ... with 6 more variables: Online <dbl>, CreditCard <dbl>, Education_1 <int>,
# Education_2 <int>, Education_3 <int>, 'Personal Loan' <dbl>
```

Part I

Data Splitting

```
#It generates the same random variables (same partition training set, same partition valid set)
set.seed(1234)

# To create a partition of 60% of our data set, we will use a caret package tool
resample = createDataPartition(mydf$Income, p=0.60, list=FALSE)

# Now, let's to create a dataframe with 60% for training sets and 40% validation sets.
train_set = mydf[resample, ]
valid_set = mydf[-resample, ]

#Now, let's do summary function to get some descriptive statistics for only income
# on our training and validation set.
summary(train_set$Income)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
8.00	39.00	64.00	73.94	98.00	205.00

```
summary(valid_set$Income)
```


Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
8.00	39.00	64.00	73.53	98.00	224.00

As we can prove, our training and validation sets are stable and well balanced.

Plotting

```
# Plotting the data would help us get a better understanding of our data subsets.

ggplot(train_set,
  aes(x=Income,y=CCAvg, color=train_set$`Personal Loan`)) + geom_point()
```



This plot show us that customers with lower income also have an average spending on credit cards lower. Or it is possible that they are likely to have less access to credit.

Normalization

Normalizing the dataset is essential to reduce the bias on the dataset.

For this purpose, we are going to use the `preProcess` function from our `Caret` package. This function uses the range method (min-max), which will transform those values using “center” (mean) and “scale” (standard deviation) as input method parameters.

To normalize the dataset, we have to exclude categorical variables such as Education, Securities Account, etc. The main variable (Personal Loan) is also excluded in this section.

Note: Normalizing in this context means transforming all those variables in smaller numbers to reduce the bias and its spreadness.

```
# Copy the original data and create new normalized dataframes
train_norm_set <- train_set
valid_norm_set <- valid_set
#traval_norm_set <- traval_set

# use preProcess() from the caret package to normalize the all the variables in our dataset.
norm_values <- preProcess(train_set[, 1:6], method=c("center", "scale"))

# Replace the columns with normalized values
train_norm_set[, 1:6] <- predict(norm_values, train_set[, 1:6])
```

```
valid_norm_set[, 1:6] <- predict(norm_values, valid_set[, 1:6])
```

```
# Now, let's see the differences between our original dataframe and our normalized dataframe.
#Calculate the descriptive stats for our normalized train set
summary(train_norm_set)
```

Age	Experience	Income	Family
Min. :-1.93925	Min. :-2.004577	Min. :-1.4330	Min. :-1.2064
1st Qu.: -0.89223	1st Qu.: -0.870011	1st Qu.: -0.7593	1st Qu.: -1.2064
Median : -0.01971	Median : 0.002733	Median : -0.2160	Median : -0.3286
Mean : 0.00000	Mean : 0.000000	Mean : 0.0000	Mean : 0.0000
3rd Qu.: 0.85282	3rd Qu.: 0.875476	3rd Qu.: 0.5229	3rd Qu.: 0.5491
Max. : 1.89984	Max. : 1.922768	Max. : 2.8483	Max. : 1.4269

CCAvg	Mortgage	Securities Account	CD Account
Min. :-1.1091	Min. :-0.5516	Min. :0.0000	Min. :0.00000
1st Qu.: -0.7054	1st Qu.: -0.5516	1st Qu.:0.0000	1st Qu.:0.00000
Median : -0.2440	Median : -0.5516	Median :0.0000	Median :0.00000
Mean : 0.0000	Mean : 0.0000	Mean :0.1099	Mean :0.06596
3rd Qu.: 0.3904	3rd Qu.: 0.4279	3rd Qu.:0.0000	3rd Qu.:0.00000
Max. : 4.6584	Max. : 5.6066	Max. :1.0000	Max. :1.00000

Online	CreditCard	Education_1	Education_2
Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.0000
1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000
Median :1.0000	Median :0.0000	Median :0.0000	Median :0.0000
Mean :0.5883	Mean :0.2965	Mean :0.4091	Mean :0.2835
3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000
Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.0000

Education_3	Personal Loan
Min. :0.0000	Min. :0.0000
1st Qu.:0.0000	1st Qu.:0.0000
Median :0.0000	Median :0.0000
Mean :0.3075	Mean :0.1023
3rd Qu.:1.0000	3rd Qu.:0.0000
Max. :1.0000	Max. :1.0000

```
#Calculate the variance for our normalized train set
var(train_norm_set[, 1:6])
```

	Age	Experience	Income	Family	CCAvg
Age	1.00000000	0.99418741	-0.04617853	-0.03281504	-0.05140346
Experience	0.99418741	1.00000000	-0.03670307	-0.03741685	-0.04765026
Income	-0.04617853	-0.03670307	1.00000000	-0.15164511	0.63670078
Family	-0.03281504	-0.03741685	-0.15164511	1.00000000	-0.09700959
CCAvg	-0.05140346	-0.04765026	0.63670078	-0.09700959	1.00000000
Mortgage	0.01965612	0.02254167	0.22577922	-0.03085638	0.10707322

	Mortgage
Age	0.01965612
Experience	0.02254167
Income	0.22577922
Family	-0.03085638
CCAvg	0.10707322
Mortgage	1.00000000

```
#Calculate the descriptive stats for out normalized valid set
summary(valid_norm_set)
```

Age	Experience	Income	Family
Min. : -1.93925	Min. : -2.004577	Min. : -1.43297	Min. : -1.20639
1st Qu.: -0.80498	1st Qu.: -0.870011	1st Qu.: -0.75928	1st Qu.: -1.20639
Median : 0.06755	Median : 0.002733	Median : -0.21597	Median : -0.32864
Mean : 0.02458	Mean : 0.029684	Mean : -0.00891	Mean : 0.04829
3rd Qu.: 0.94007	3rd Qu.: 0.875476	3rd Qu.: 0.52292	3rd Qu.: 1.42685
Max. : 1.89984	Max. : 2.010042	Max. : 3.26116	Max. : 1.42685

CCAvg	Mortgage	Securities Account	CD Account
Min. : -1.10913	Min. : -0.551650	Min. : 0.0000	Min. : 0.00000
1st Qu.: -0.70540	1st Qu.: -0.551650	1st Qu.: 0.0000	1st Qu.: 0.00000
Median : -0.18632	Median : -0.551650	Median : 0.0000	Median : 0.00000
Mean : 0.02148	Mean : -0.009318	Mean : 0.0961	Mean : 0.05205
3rd Qu.: 0.33275	3rd Qu.: 0.427850	3rd Qu.: 0.0000	3rd Qu.: 0.00000
Max. : 4.65842	Max. : 5.432028	Max. : 1.0000	Max. : 1.00000

Online	CreditCard	Education_1	Education_2
Min. : 0.0000	Min. : 0.0000	Min. : 0.0000	Min. : 0.0000
1st Qu.: 0.0000	1st Qu.: 0.0000	1st Qu.: 0.0000	1st Qu.: 0.0000
Median : 1.0000	Median : 0.0000	Median : 0.0000	Median : 0.0000
Mean : 0.6096	Mean : 0.2903	Mean : 0.4344	Mean : 0.2763
3rd Qu.: 1.0000	3rd Qu.: 1.0000	3rd Qu.: 1.0000	3rd Qu.: 1.0000
Max. : 1.0000	Max. : 1.0000	Max. : 1.0000	Max. : 1.0000

Education_3	Personal Loan
Min. : 0.0000	Min. : 0.00000
1st Qu.: 0.0000	1st Qu.: 0.00000
Median : 0.0000	Median : 0.00000
Mean : 0.2893	Mean : 0.08659
3rd Qu.: 1.0000	3rd Qu.: 0.00000
Max. : 1.0000	Max. : 1.00000

```
#Calculate the variance for out normalized valid set
var(valid_norm_set[, 1:6])
```

	Age	Experience	Income	Family	CCAvg
Age	1.00107163	0.99691546	-0.06890733	-0.06847225	-0.05433237
Experience	0.99691546	1.00426859	-0.06142022	-0.07729387	-0.05523825
Income	-0.06890733	-0.06142022	1.00250134	-0.16918536	0.67392892
Family	-0.06847225	-0.07729387	-0.16918536	1.03608725	-0.13259335
CCAvg	-0.05433237	-0.05523825	0.67392892	-0.13259335	1.04029149
Mortgage	-0.06036765	-0.05985953	0.17152809	-0.00421576	0.11276287

	Mortgage
Age	-0.06036765
Experience	-0.05985953
Income	0.17152809
Family	-0.00421576
CCAvg	0.11276287
Mortgage	0.93293385

As we can see, after normalizing the training set, the mean and variance of the columns are between 0 and 1. It proves that the data is in equal scale to be analyzed. Regarding the validation set, it is not true. We will use the mean and standard deviation from the training set to normalize the validation and testing set.

Training k-NN Model

Now, let apply the k-NN model where $k = 1$. First, create the x variables (all the ones we want to train) and create another one to save the Y variable, which is Personal Loan (last column). This column has values of 1=Yes, or 0=No.

We will use “drop = TRUE” argument to transform the dataframe into a vector because k-NN works only with vectors.

```
#First, create X dataframe and Y vector
train_predictors<-train_norm_set[,1:13, drop = TRUE]
valid_predictors<-valid_norm_set[,1:13, drop = TRUE]
train_labels <-train_norm_set[,14, drop = TRUE]
valid_labels  <-valid_norm_set[,14, drop = TRUE]

#Run the model using k = 1
set.seed(1234)
my_knn <-knn(train_predictors,
              valid_predictors,
              cl=train_labels,
              k=1 )

# See the 6 first values of predicted class in the validation set
head(my_knn)
```

```
[1] 0 0 0 0 0 0
Levels: 0 1
```

```
# To summarized the model
summary(my_knn)
```

```
      0      1
1869  129
```

Confusion Matrix

```
# Create a confusion matrix
conf_matrix <- CrossTable(x=valid_labels,y=my_knn, prop.chisq = FALSE)
```

```
      Cell Contents
|-----|
|              N |
|      N / Row Total |
|      N / Col Total |
|      N / Table Total |
|-----|
```

```
Total Observations in Table:  1998
```

valid_labels	my_knn		Row Total
	0	1	
0	1807	18	1825
	0.990	0.010	0.913
	0.967	0.140	
	0.904	0.009	
1	62	111	173
	0.358	0.642	0.087
	0.033	0.860	
	0.031	0.056	
Column Total	1869	129	1998
	0.935	0.065	

The confusion matrix show us that the model wrongly predicted 80 customers. The model is classified 95.99% correctly.

Probability Output

```
# Create a new variable for our probability
set.seed(1234)
my_knnprob <- knn(train_predictors,
  valid_predictors,
  cl=train_labels, k=1, prob=TRUE )

class_prob<-attr(my_knnprob, 'prob')

# See the first rows
head(class_prob)
```

```
[1] 1 1 1 1 1 1
```

Calcutale the accuracy, recall, precision, specificity

```
#Calcutale the accuracy
k1_accuracy <- (conf_matrix$t[2,2] + conf_matrix$t[1,1]) / sum(conf_matrix$t)
print(k1_accuracy)
```

```
[1] 0.95996
```

```
#Calcutale the recall
k1_recall <- conf_matrix$t[2,2]/ (conf_matrix$t[2,2] + conf_matrix$t[2,1])
print(k1_recall)
```

```
[1] 0.6416185
```

```
#Calcutale the precision
k1_precision <- conf_matrix$t[2,2]/ (conf_matrix$t[2,2] + conf_matrix$t[1,2])
print(k1_precision)
```

```
[1] 0.8604651
```

```
#Calcutale the specificity
k1_specificity <- conf_matrix$t[1,1]/ (conf_matrix$t[1,1] + conf_matrix$t[1,2])
print(k1_specificity)
```

```
[1] 0.990137
```

As we can determine, the model is learning well.

Create a new observation

Now, let's add a new customer and run the k-NN model the validation set.

```
# Create a new observation
new_obs <- c(40, 10, 84, 2, 2, 0, 0, 0, 1, 1, 0, 1, 0)
```

Normalize the new observation

```
# To select only the first six columns that we want to normalize
x <- as.data.frame(t(new_obs[1:6]))
colnames(x) <- c('Age', 'Expereince', 'Income', 'Family', 'CCAvg', 'Mortgage')

#To normalize my new obervation
n=(x-norm_values$mean)/norm_values$std

# To get the same dimensions
m = c(n$Age,n$Expereince,n$Income,n$Family,n$CCAvg,n$Mortgage,new_obs[7], new_obs[8],
      new_obs[9], new_obs[10], new_obs[11], new_obs[12], new_obs[13])

# To show how my the variable looks like
m
```

```
[1] -0.45596688 -0.87001050  0.21866737 -0.32864440  0.04437672 -0.55164997
[7]  0.00000000  0.00000000  1.00000000  1.00000000  0.00000000  1.00000000
[13]  0.00000000
```

Run K-NN Model with the new observation

```

# Calculate the probability for the new observation
my_knnprob2 <-knn(train_predictors,
                  m,
                  cl=train_labels,
                  k=1, prob = TRUE)

# To get the probability output
class_prob2<-attr(my_knnprob2, 'prob')
# Show the first lines
head(class_prob2)

```

```
[1] 1
```

Explain how would this customer be classified?

The probability output the new customer “1.” It means that the new customer will accept the personal loan offer from the Universal Bank.

Determining the optimal using Hyperparameter Tuning

Let's find the optimal k using tuning parameters

```

set.seed(1234)
Search_grid <- expand.grid(k=c(1:20))
train_predict_labels <- train_predictors
train_predict_labels$Personal_Loan = train_labels

modeltest<-train(factor(Personal_Loan)~Age+Experience+Income+Family+
                   CCAvg+Mortgage+'Securities Account'+ 'CD Account'+Online+
                   CreditCard+Education_1+Education_2+Education_3,
                   data = train_predict_labels, method="knn",
                   tuneGrid=Search_grid,
                   preProcess='range')

# To show the result
modeltest

```

k-Nearest Neighbors

```

3002 samples
 13 predictor
 2 classes: '0', '1'

```

Pre-processing: re-scaling to [0, 1] (13)

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 3002, 3002, 3002, 3002, 3002, 3002, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
1	0.9536046	0.7263191
2	0.9471997	0.6821014
3	0.9458861	0.6670179

4	0.9452363	0.6582538
5	0.9446240	0.6448406
6	0.9441927	0.6366618
7	0.9432076	0.6260626
8	0.9411801	0.6052238
9	0.9399465	0.5922334
10	0.9381334	0.5738815
11	0.9370963	0.5640625
12	0.9356435	0.5523524
13	0.9350647	0.5445418
14	0.9330659	0.5249848
15	0.9319625	0.5132661
16	0.9313246	0.5060602
17	0.9298700	0.4910055
18	0.9283937	0.4769321
19	0.9273087	0.4643694
20	0.9261396	0.4539903

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was $k = 1$.

As we can see, the hypertuning parameter uses a resampling of Bootstrapped (25 reps), which is very robust, and it gives the optimal $k = 1$

Discuss the choice of k that balances between overfitting and ignoring the predictor information?

Overfitting is when the model is capturing the personal characteristics from the training set. We can prove that comparing the accuracy from the training and validation set.

Based on our outputs showing the accuracy of the model trained on different data partitioning, we can prove out the model is not overfitting because those values have similar results.

Show the confusion matrix for the validation data that results from using the best k and explain different error types that you observe.

Confusion Matrix using the optimal k .

In the universal Bank problem, the hypertuning parameter uses a resampling of Bootstrapped (25 reps), which is very robust, and gives the optimal $k = 1$. So, we will continue using $k=1$.

Part II

Data Splitting for the second part

Here we are going to divide the dataframe in three sections: Training, Validation, and Testing

```
#It generates the same random variables (same partition training set, same partition valid set)
set.seed(1234)

# To create a partition of 80% of our data set, we will use a caret package tool
resample2 = createDataPartition(mydf$Income, p=0.80, list=FALSE)

# Now, let's to create a dataframe with 50% for training sets, 30% validation sets,
# and 20% testing sets.
```

```

traval_set2 = mydf[resample2, ] # It will get 80%
test_set2 = mydf[-resample2, ] # It will get 20%

# To create a partition of 50% of our data set, we will use a caret package tool
resample3 = createDataPartition(traval_set2$Income, p=0.80, list=FALSE)

# To share the data form training and validation
train_set2 = mydf[resample3, ] # It will get 50%
valid_set2 = mydf[-resample3, ] # It will get 30%

#Now, let's do summary function to get some descriptive statistics for only income
# on our training and validation set.
summary(train_set2$Income)

```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
8.00	39.00	63.00	73.57	98.00	224.00

```
summary(valid_set2$Income)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
8.00	39.00	64.00	74.14	99.00	218.00

```
summary(test_set2$Income)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
8.00	39.00	63.00	73.33	98.00	224.00

Normalizing step

This steps is to renormalize our dataframes. It is important to notice that we have to renormalize our traval (training+validation) set.

```

# Copy the original data and create new normalization dataframes
train_norm_set1 <- train_set2
valid_norm_set1 <- valid_set2
traval_norm_set1 <- traval_set2
test_norm_set1 <- test_set2

# use preProcess() from the caret package to normalize the all the variables for the training set.
norm_values2 <- preProcess(train_set2[, 1:6], method=c("center", "scale"))

# Replace the columns with normalized values for the training and validation set
train_norm_set1[, 1:6] <- predict(norm_values2, train_set2[, 1:6])
valid_norm_set1[, 1:6] <- predict(norm_values2, valid_set2[, 1:6])

# use preProcess() from the caret package to normalize the all the variables for the training and valid
norm_values3 <- preProcess(traval_set2[, 1:6], method=c("center", "scale"))

# Replace the columns with normalized values
traval_norm_set1[, 1:6] <- predict(norm_values3, traval_set2[, 1:6])

```

```
test_norm_set1[, 1:6] <- predict(norm_values3, test_set2[, 1:6])
```

```
# Now, let's see the differences between our original dataframe and our normalized dataframe.
summary(train_norm_set1)
```

Age	Experience	Income	Family
Min. :-1.93922	Min. :-2.00430	Min. :-1.4372	Min. :-1.2185
1st Qu.: -0.89920	1st Qu.: -0.87802	1st Qu.: -0.7577	1st Qu.: -1.2185
Median : -0.03252	Median : -0.01166	Median : -0.2317	Median : -0.3545
Mean : 0.00000	Mean : 0.00000	Mean : 0.0000	Mean : 0.0000
3rd Qu.: 0.83416	3rd Qu.: 0.85471	3rd Qu.: 0.5355	3rd Qu.: 1.3736
Max. : 1.87417	Max. : 1.98098	Max. : 3.2971	Max. : 1.3736

CCAvg	Mortgage	Securities Account	CD Account
Min. :-1.1112	Min. :-0.5616	Min. :0.0000	Min. :0.00000
1st Qu.: -0.7104	1st Qu.: -0.5616	1st Qu.:0.0000	1st Qu.:0.00000
Median : -0.2522	Median : -0.5616	Median :0.0000	Median :0.00000
Mean : 0.0000	Mean : 0.0000	Mean :0.1074	Mean :0.05963
3rd Qu.: 0.3204	3rd Qu.: 0.4442	3rd Qu.:0.0000	3rd Qu.:0.00000
Max. : 4.6152	Max. : 5.5223	Max. :1.0000	Max. :1.00000

Online	CreditCard	Education_1	Education_2
Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.000
1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.000
Median :1.0000	Median :0.0000	Median :0.0000	Median :0.000
Mean :0.5982	Mean :0.2882	Mean :0.4131	Mean :0.285
3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.000
Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.000

Education_3	Personal Loan
Min. :0.0000	Min. :0.0000
1st Qu.:0.0000	1st Qu.:0.0000
Median :0.0000	Median :0.0000
Mean :0.3019	Mean :0.1005
3rd Qu.:1.0000	3rd Qu.:0.0000
Max. :1.0000	Max. :1.0000

```
#Calculate the variance
var(train_norm_set1[, 1:6])
```

	Age	Experience	Income	Family	CCAvg
Age	1.00000000	0.99429778	-0.04732990	-0.05657548	-0.03250474
Experience	0.99429778	1.00000000	-0.03896261	-0.06246274	-0.03155934
Income	-0.04732990	-0.03896261	1.00000000	-0.14804842	0.63761632
Family	-0.05657548	-0.06246274	-0.14804842	1.00000000	-0.09710343
CCAvg	-0.03250474	-0.03155934	0.63761632	-0.09710343	1.00000000
Mortgage	-0.02014256	-0.01665128	0.21263362	-0.02366249	0.11644228

	Mortgage
Age	-0.02014256
Experience	-0.01665128
Income	0.21263362
Family	-0.02366249
CCAvg	0.11644228
Mortgage	1.00000000

```
summary(valid_norm_set1)
```

Age	Experience	Income	Family
Min. : -1.939220	Min. : -2.004297	Min. : -1.43715	Min. : -1.21851
1st Qu.: -0.812536	1st Qu.: -0.878023	1st Qu.: -0.75770	1st Qu.: -1.21851
Median : -0.032524	Median : -0.011658	Median : -0.20975	Median : -0.35447
Mean : -0.008892	Mean : -0.007222	Mean : 0.01247	Mean : -0.03327
3rd Qu.: 0.834156	3rd Qu.: 0.854707	3rd Qu.: 0.55738	3rd Qu.: 0.50958
Max. : 1.874171	Max. : 1.980981	Max. : 3.16561	Max. : 1.37362

CCAvg	Mortgage	Securities Account	CD Account
Min. : -1.111203	Min. : -0.56158	Min. : 0.00000	Min. : 0.00000
1st Qu.: -0.727536	1st Qu.: -0.56158	1st Qu.: 0.00000	1st Qu.: 0.00000
Median : -0.252247	Median : -0.56158	Median : 0.00000	Median : 0.00000
Mean : -0.004081	Mean : -0.01246	Mean : 0.09905	Mean : 0.06177
3rd Qu.: 0.377655	3rd Qu.: 0.40474	3rd Qu.: 0.00000	3rd Qu.: 0.00000
Max. : 4.615173	Max. : 5.69979	Max. : 1.00000	Max. : 1.00000

Online	CreditCard	Education_1	Education_2
Min. : 0.0000	Min. : 0.0000	Min. : 0.0000	Min. : 0.0000
1st Qu.: 0.0000	1st Qu.: 0.0000	1st Qu.: 0.0000	1st Qu.: 0.0000
Median : 1.0000	Median : 0.0000	Median : 0.0000	Median : 0.0000
Mean : 0.5943	Mean : 0.3044	Mean : 0.4302	Mean : 0.2727
3rd Qu.: 1.0000	3rd Qu.: 1.0000	3rd Qu.: 1.0000	3rd Qu.: 1.0000
Max. : 1.0000	Max. : 1.0000	Max. : 1.0000	Max. : 1.0000

Education_3	Personal Loan
Min. : 0.0000	Min. : 0.00000
1st Qu.: 0.0000	1st Qu.: 0.00000
Median : 0.0000	Median : 0.00000
Mean : 0.2972	Mean : 0.08792
3rd Qu.: 1.0000	3rd Qu.: 0.00000
Max. : 1.0000	Max. : 1.00000

```
#Calculate the variance
var(valid_norm_set1[, 1:6])
```

	Age	Experience	Income	Family	CCAvg
Age	0.964380417	0.9588103166	-0.06975031	-0.02660781	-0.08601218
Experience	0.958810317	0.9646945756	-0.06043017	-0.03293567	-0.08234398
Income	-0.069750310	-0.0604301674	1.05058126	-0.17440245	0.67881739
Family	-0.026607815	-0.0329356683	-0.17440245	0.95345056	-0.12880977
CCAvg	-0.086012183	-0.0823439846	0.67881739	-0.12880977	1.00486295
Mortgage	0.001065446	0.0002804293	0.20349936	-0.01467524	0.09941529

	Mortgage
Age	0.0010654464
Experience	0.0002804293
Income	0.2034993599
Family	-0.0146752358
CCAvg	0.0994152864
Mortgage	1.0168521539

Training k-NN Model for Validation Set

Training the k-NN Model using the training and validation set


```

#First, create the variables for The Y variable which is Personal Loan (last column).
train_predictors1 <-train_norm_set1[,1:13, drop = TRUE]
valid_predictors1 <-valid_norm_set1[,1:13, drop = TRUE]
train_labels1 <-train_norm_set1[,14, drop = TRUE]
valid_labels1 <-valid_norm_set1[,14, drop = TRUE]

#Run the model using k = 1
set.seed(1234)
my_knn2 <-knn(train_predictors1,
               valid_predictors1,
               cl=train_labels1,
               k=1 )

# See the 6 first values of predicted class in the validation set
head(my_knn2)

```

```

[1] 0 0 0 1 0 0
Levels: 0 1

```

```

# To summarized the model
summary(my_knn2)

```

```

      0      1
1664  133

```

Determining the optimal using Hyperparameter Tuning for Validation Set

Let's find the optimal k using tuning parameters

```

set.seed(1234)
Search_grid <- expand_grid(k=c(1:20))
train_predict_labels1 <- train_predictors1
train_predict_labels1$Personal_Loan = train_labels1
modeltest1<-train(factor(Personal_Loan)~Age+Experience+Income+Family+
                    CCAvg+Mortgage+'Securities Account'+ 'CD Account'+Online+
                    CreditCard+Education_1+Education_2+Education_3,
                    data = train_predict_labels1, method="knn",
                    tuneGrid=Search_grid,
                    preProcess='range')

# To show the result
modeltest1

```

k-Nearest Neighbors

```

3203 samples
  13 predictor
  2 classes: '0', '1'

```

```

Pre-processing: re-scaling to [0, 1] (13)
Resampling: Bootstrapped (25 reps)

```

Summary of sample sizes: 3203, 3203, 3203, 3203, 3203, 3203, ...
Resampling results across tuning parameters:

k	Accuracy	Kappa
1	0.9503554	0.6930832
2	0.9452988	0.6577146
3	0.9438377	0.6433481
4	0.9435997	0.6338016
5	0.9429516	0.6196916
6	0.9424734	0.6102541
7	0.9416765	0.5987671
8	0.9402743	0.5819964
9	0.9384071	0.5633323
10	0.9367152	0.5464649
11	0.9360400	0.5387656
12	0.9342771	0.5207271
13	0.9333490	0.5121482
14	0.9319518	0.4969679
15	0.9309023	0.4855445
16	0.9297518	0.4742680
17	0.9284311	0.4594290
18	0.9271093	0.4455629
19	0.9261323	0.4325538
20	0.9248099	0.4174314

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was k = 1.

Training k-NN Model for Test Set

Training the k-NN Model using the training set (training + validation) and testing set

```
#First, create the variables for The Y variable which is Personal Loan (last column).
traval_predictors1 <-traval_norm_set1[,1:13, drop = TRUE]
test_predictors1 <-test_norm_set1[,1:13, drop = TRUE]
traval_labels1 <-traval_norm_set1[,14, drop = TRUE]
test_labels1 <-test_norm_set1[,14, drop = TRUE]
```

```
#Run the model using k = 1
```

```
set.seed(1234)
my_knn3 <-knn(traval_predictors1,
              test_predictors1,
              cl=traval_labels1,
              k=1 )
```

```
# See the 6 first values of predicted class in the validation set
```

```
head(my_knn3)
```

```
[1] 0 0 0 0 1 1
Levels: 0 1
```

```
# To summarized the model
summary(my_knn3)
```

```
0    1
914 84
```

Determining the optimal using Hyperparameter Tuning for Test Set

Let's find the optimal k using tuning parameters

```
set.seed(1234)
Search_grid <- expand.grid(k=c(1:20))
traval_predict_labels1 <- traval_predictors1
traval_predict_labels1$Personal_Loan = traval_labels1
modeltest2 <- train(factor(Personal_Loan)~Age+Experience+Income+Family+
                     CCAvg+Mortgage+'Securities Account'+ 'CD Account'+Online+
                     CreditCard+Education_1+Education_2+Education_3,
                     data = traval_predict_labels1, method="knn",
                     tuneGrid=Search_grid,
                     preProcess='range')

# To show the result
modeltest2
```

k-Nearest Neighbors

```
4002 samples
 13 predictor
 2 classes: '0', '1'
```

Pre-processing: re-scaling to [0, 1] (13)

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 4002, 4002, 4002, 4002, 4002, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
1	0.9591370	0.7342128
2	0.9548035	0.7010894
3	0.9536538	0.6864198
4	0.9530212	0.6769620
5	0.9521833	0.6646934
6	0.9516948	0.6577040
7	0.9509923	0.6461417
8	0.9498775	0.6347656
9	0.9485197	0.6197664
10	0.9475126	0.6085382
11	0.9466458	0.5993160
12	0.9446654	0.5787931
13	0.9436382	0.5686613
14	0.9427109	0.5591373
15	0.9415228	0.5469459
16	0.9410365	0.5419454

```

17  0.9394884  0.5257653
18  0.9384533  0.5147062
19  0.9372660  0.5024011
20  0.9358777  0.4867357

```

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was $k = 1$.

K-NN for Training set

```

#Run the model using k = 1
set.seed(1234)
my_knn0 <- knn(train_predictors1,
               train_predictors1,
               cl=train_labels1,
               k=1 )

# See the 6 first values of predicted class in the validation set
head(my_knn0)

```

```

[1] 0 0 0 0 0 0
Levels: 0 1

```

```

# To summarized the model
summary(my_knn0)

```

```

      0      1
2881  322

```

Confusion Matrix for Training set

```

# Create a confusion matrix
conf_matrix0 <- CrossTable(x=train_labels1,y=my_knn0, prop.chisq = FALSE)

```

```

      Cell Contents
|-----|
|              N |
|      N / Row Total |
|      N / Col Total |
|      N / Table Total |
|-----|

```

Total Observations in Table: 3203

train_labels1	my_knn0		Row Total
	0	1	
0	2881	0	2881
	1.000	0.000	0.899
	1.000	0.000	
	0.899	0.000	
1	0	322	322
	0.000	1.000	0.101
	0.000	1.000	
	0.000	0.101	
Column Total	2881	322	3203
	0.899	0.101	

This confusion matrix is showing a perfect output because it is classifying the same data that is used to train the model.

I showed this confusion matrix just to confirm its accuracy. However, it is not fair to compare this result against the confusion matrix obtained from the validation and test set.

Calculate the accuracy, recall, precision, specificity for Training set

```
#Calculate the accuracy
k1_accuracy0 <- (conf_matrix0$t[2,2] + conf_matrix0$t[1,1]) / sum(conf_matrix0$t)
print(k1_accuracy0)
```

```
[1] 1
```

```
#Calculate the recall
k1_recall0 <- conf_matrix0$t[2,2] / (conf_matrix0$t[2,2] + conf_matrix0$t[2,1])
print(k1_recall0)
```

```
[1] 1
```

```
#Calculate the precision
k1_precision0 <- conf_matrix0$t[2,2] / (conf_matrix0$t[2,2] + conf_matrix0$t[1,2])
print(k1_precision0)
```

```
[1] 1
```

```
#Calculate the specificity
k1_specificity0 <- conf_matrix0$t[1,1] / (conf_matrix0$t[1,1] + conf_matrix0$t[1,2])
print(k1_specificity0)
```

```
[1] 1
```

Confusion Matrix for Validation set

```
# Create a confusion matrix
conf_matrix1 <- CrossTable(x=valid_labels1,y=my_knn2, prop.chisq = FALSE)
```

```
      Cell Contents
|-----|
|              N |
|      N / Row Total |
|      N / Col Total |
|      N / Table Total |
|-----|
```

Total Observations in Table: 1797

	my_knn2		
valid_labels1	0	1	Row Total
0	1622	17	1639
	0.990	0.010	0.912
	0.975	0.128	
	0.903	0.009	
1	42	116	158
	0.266	0.734	0.088
	0.025	0.872	
	0.023	0.065	
Column Total	1664	133	1797
	0.926	0.074	

This confusion matrix shows us that utilizing the validation set the model is wrongly classifying 59 customers, which means 21 extra people were correctly classified.

Probability Output for Validation set

```
# Create a new variable for our probability
set.seed(1234)
my_knnprob2 <- knn(train_predictors1,
  valid_predictors1,
  cl=train_labels1, k=1, prob=TRUE )

class_prob2 <- attr(my_knnprob2, 'prob')
```

```
# See the first rows
head(class_prob2)
```

```
[1] 1 1 1 1 1 1
```

Calcutale the accuracy, recall, precision, specificity for Validation set

```
#Calcutale the accuracy
k1_accuracy2 <- (conf_matrix1$t[2,2] + conf_matrix1$t[1,1]) / sum(conf_matrix1$t)
print(k1_accuracy2)
```

```
[1] 0.9671675
```

```
#Calcutale the recall
k1_recall2 <- conf_matrix1$t[2,2] / (conf_matrix1$t[2,2] + conf_matrix1$t[2,1])
print(k1_recall2)
```

```
[1] 0.7341772
```

```
#Calcutale the precision
k1_precision2 <- conf_matrix1$t[2,2] / (conf_matrix1$t[2,2] + conf_matrix1$t[1,2])
print(k1_precision2)
```

```
[1] 0.8721805
```

```
#Calcutale the specificity
k1_specificity2 <- conf_matrix1$t[1,1] / (conf_matrix1$t[1,1] + conf_matrix1$t[1,2])
print(k1_specificity2)
```

```
[1] 0.9896278
```

Confusion Matrix for Testing set

```
# Create a confusion matrix
conf_matrix2 <- CrossTable(x=test_labels1,y=my_knn3, prop.chisq = FALSE)
```

```

      Cell Contents
|-----|
|              N |
|      N / Row Total |
|      N / Col Total |
|      N / Table Total |
|-----|
```

Total Observations in Table: 998

test_labels1	my_knn3		Row Total
	0	1	
0	884	9	893
	0.990	0.010	0.895
	0.967	0.107	
	0.886	0.009	
1	30	75	105
	0.286	0.714	0.105
	0.033	0.893	
	0.030	0.075	
Column Total	914	84	998
	0.916	0.084	

When applying the testing, the model is wrongly classifying 39 customers. It proves that the model is performing better than the previous one.

Probability Output for Testing set

```
# Create a new variable for our probability
set.seed(1234)
my_knnprob3 <- knn(traval_predictors1,
  test_predictors1,
  cl=traval_labels1, k=1, prob=TRUE )

class_prob3<-attr(my_knnprob3, 'prob')

# See the first rows
head(class_prob3)
```

```
[1] 1 1 1 1 1 1
```

Calculate the accuracy, recall, precision, specificity for Testing set

```
#Calculate the accuracy
k1_accuracy3 <- (conf_matrix2$t[2,2] + conf_matrix2$t[1,1]) / sum(conf_matrix2$t)
print(k1_accuracy3)
```

```
## [1] 0.9609218
```



```
#Calcutale the recall
k1_recall3 <- conf_matrix2$t[2,2] / (conf_matrix2$t[2,2] + conf_matrix2$t[2,1])
print(k1_recall3)
```

```
## [1] 0.7142857
```

```
#Calcutale the precision
k1_precision3 <- conf_matrix2$t[2,2] / (conf_matrix2$t[2,2] + conf_matrix2$t[1,2])
print(k1_precision3)
```

```
## [1] 0.8928571
```

```
#Calcutale the specificity
k1_specificity3 <- conf_matrix2$t[1,1] / (conf_matrix2$t[1,1] + conf_matrix2$t[1,2])
print(k1_specificity3)
```

```
## [1] 0.9899216
```

Run K-NN Model to try to predict again the same customer with the new sets

```
# Calculate the probability for the new observation
my_finalpred <-knn(test_predictors1,
                  m,
                  cl=test_labels1,
                  k=1, prob = TRUE)

# To get the probability output
class_finalprob<-attr(my_finalpred , 'prob')
# Show the first lines
head(class_finalprob)
```

```
[1] 1
```

In conclusion, the new customer is going to be classify as accepting the personal loan form the Universal Bank from the new marketing campaing.